



Kubernetes

Verizon Study Groups



Agenda

- **Introduction**
- **Prerequisites**
- **What is Kubernetes?**
- **Why you need Kubernetes?**
- **Kubernetes Components**
 - **Control Plane Components**
 - **Node Components**



Prerequisites

- Docker
- Setup lab environment for hands-on (kubectl, minikube,)
<https://kubernetes.io/docs/tasks/tools/>

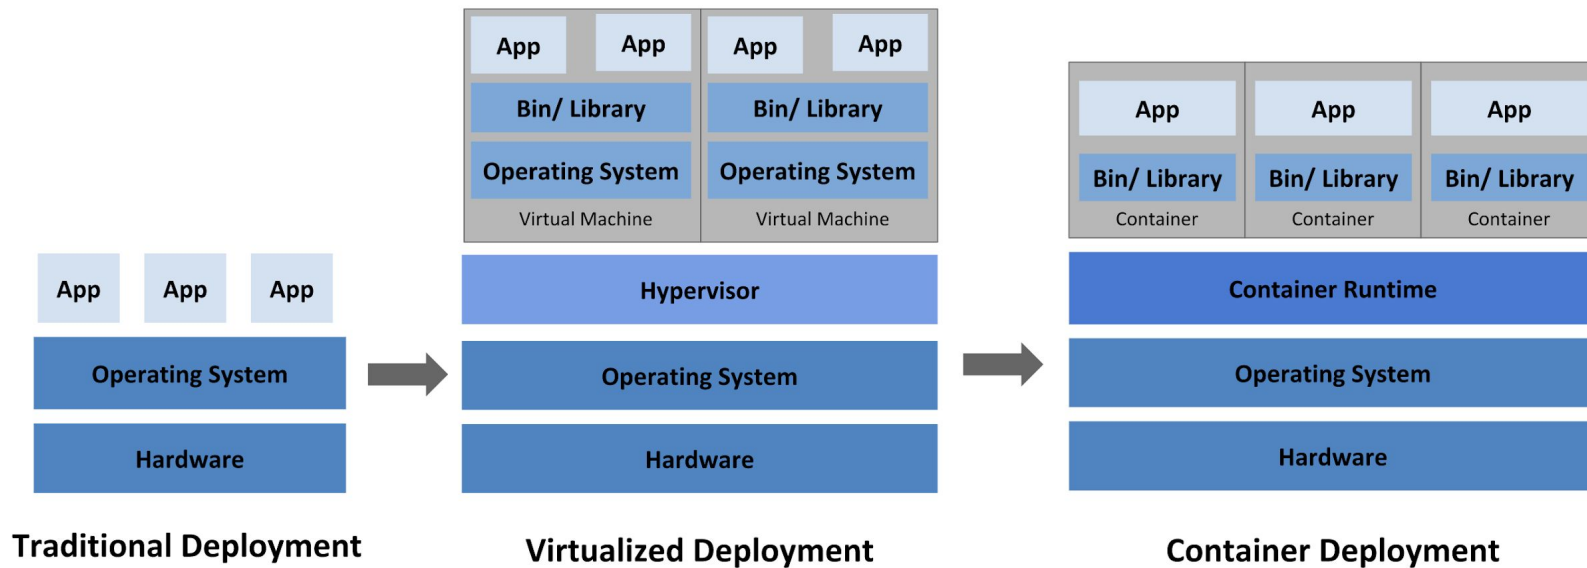
[Verizon Kubernetes Sandbox/Playground](#)

What is Kubernetes?

Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation.



Evolution





Why you need Kubernetes?

Kubernetes provides you with a framework to run distributed systems resiliently. It takes care of scaling and failover for your application, provides deployment patterns, and more



Kubernetes Components

https://www.youtube.com/watch?v=8C_SCDbUJTg

Kubernetes Architecture



Master

Manage, Plan, Schedule, Monitor
Nodes

ETCD
CLUSTER

kube-
apiserver

Kube
Controller
Manager

kube-scheduler



Worker Nodes

Host Application as Containers

kubelet

Kube-proxy

Container Runtime Engine

Run containers



kubelet

Kube-proxy

Container Runtime Engine

Run containers





kube-apiserver

The core of Kubernetes' [control plane](#)

The API server exposes an HTTP API which can be used to query the cluster



ETCD - <https://etcd.io/>

etcd is a consistent and highly-available key value store used as Kubernetes' backing store for all cluster data.

KEY	VALUE
Name	Smith
Age	33
Location	New York

```
{  
  Name : Smith,  
  Age   : 33,  
  Location : New York  
}
```



kube-scheduler

Control plane component that watches for newly created Pods with no assigned node, and selects a node for them to run on.



kube-controller-manager

Control Plane component that runs [controller](#) processes.



cloud-controller-manager

A Kubernetes control plane component that embeds cloud-specific control logic



References

<https://kubernetes.io/docs/home/>

https://www.youtube.com/watch?v=8C_SCDbUJTg

<https://etcd.io/>



Agenda - Day 2

- **Pods**
- **Deployments (Replicasets)**
- **Services**
- **Namespaces**
- **YAML**
- **Static Pods**
- **DaemonSets**

Container

Baby computer inside a computer

Server

Container

- Alpine
- Python
- Setup

APP

```
FROM python:alpine
```

```
WORKDIR /app
```



```
RUN execute_cmds
```

```
EXPOSE 80
```

```
CMD ["python", "app.py"]
```




**Where should I live?
How to talk to other containers?
How to talk to the world?
What happens if I get sick?**









POD

10.1.2.3




10.1.2.4






Deployments






Services

service1.corp.xxx.com



service2.corp.xxx.com





YAML (Yet Another Markup Language)

```
apiVersion:
```

```
kind:
```

```
metadata:
```

```
spec:
```



Namespaces

Virtual cluster inside your Kubernetes cluster

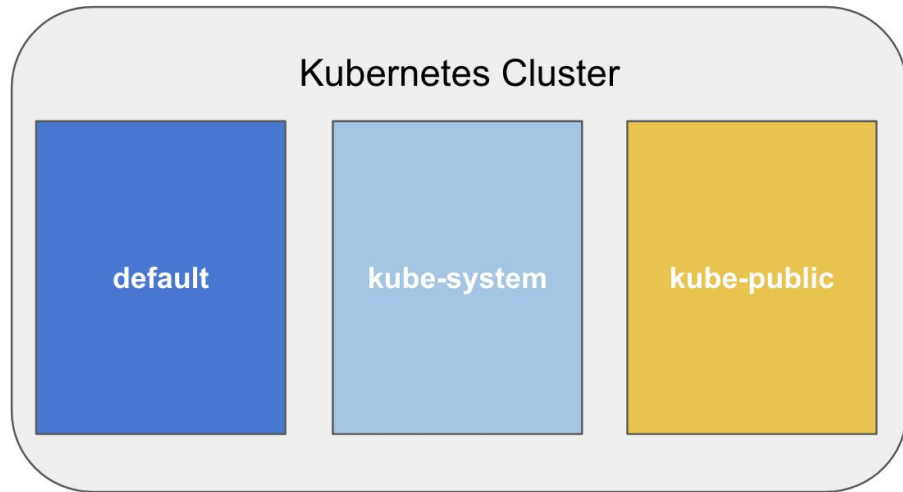
//creating namespace with single command

```
$ kubectl create namespace dev
```

namespace/dev created.



Imperative Way





PODS

```
kubectl run nginx --image=nginx
```

```
kubectl create -f nginx.yaml
```

<https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands#run>



Deployments

```
kubectl create deployment nginx --image nginx
```

```
kubectl create -f nginx.yaml
```

<https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands#-em-deployment-em->



Services

```
kubectl create service
```

```
kubectl expose pod valid-pod --port=444 --name=front
```

```
kubectl create -f nginx.yaml
```

<https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands#expose>



Static Pods

Static Pods are managed directly by the **kubelet daemon** on a specific node, without the API server observing them. Unlike Pods that are managed by the control plane (for example, a Deployment); instead, the kubelet watches each static Pod (and restarts it if it fails).



Multiple Pods

Resources that manage one or more Pods:

- Deployment
- StatefulSet
- DaemonSet



DaemonSets

A *DaemonSet* ensures that all (or some) Nodes run **a copy** of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected. Deleting a DaemonSet will clean up the Pods it created.

Use Cases:

- Logs, monitoring daemons



References

YAML: <https://www.youtube.com/watch?v=1uFVr15xDGg>
<https://www.youtube.com/watch?v=qmDzcu5uY1I>



Lab Practice

- List all the namespaces
- List all the pods in your namespace
- Create a namespace.
- Create a pod with image nginx in the namespace created
- Create a service to expose the pod on port 8080
- Create a deployment with “redis” image on port 6379 and expose a service.
- Increase the number of replicas to 3



Agenda - Day 3

- **Services - Kinds**
- **Scheduling**
- **Labels and Selectors**
- **Taints and Tolerations**
- **Node Selectors**
- **Node Affinity**
- **Lab Practice**



Services - Kinds

Kubernetes ServiceTypes allow you to specify what kind of Service you want.


Type values:

- ClusterIP (default)
- NodePort
- LoadBalancer
- ExternalName


Services Types



NodePort



ClusterIP



LoadBalancer



Services - Reference

<https://www.youtube.com/watch?v=5lzUpDtmWgM>

<https://www.youtube.com/watch?v=T4Z7visMM4E>



Scheduling

What to schedule?

Where to schedule?

Manual Scheduling - Use nodeName before creating the pod

<https://kubernetes.io/docs/concepts/scheduling-eviction/kube-scheduler/>



Scheduling - Reference

<https://www.youtube.com/watch?v=rDCWxkvPIAw>

<https://www.youtube.com/watch?v=E3ExWruji7g>



Labels and Selectors

Labels - Grouping

Selectors - Filtering

```
"metadata": {  
  "labels": {  
    "key1": "value1",  
    "key2": "value2"  
  }  
}
```



Labels and Selectors

Example labels:

- `"release" : "stable", "release" : "canary"`
- `"environment" : "dev", "environment" : "qa", "environment" : "production"`
- `"tier" : "frontend", "tier" : "backend", "tier" : "cache"`
- `"partition" : "customerA", "partition" : "customerB"`
- `"track" : "daily", "track" : "weekly"`



apiVersion: v1

kind: Pod

metadata:

name: label-demo

labels:

environment: production

app: nginx

spec:

containers:

- **name:** nginx

image: nginx:1.14.2

ports:

- **containerPort:** 80



Selector Examples

```
kubectl get pods --show-labels
```

```
kubectl get pods -l environment=production,tier=frontend
```

```
kubectl get pods -l 'environment in (production),tier in (frontend)'
```

```
kubectl get pods -l 'environment in (production, qa)'
```

```
kubectl get pods -l 'environment,environment notin (frontend)'
```

```
kubectl get pods --selector version=1.0
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
```

```
  labels:
```

```
    app: my-first
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      app: nginx
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: nginx
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx
```

```
          image: nginx:1.14.2
```

```
          ports:
```

```
            containerPort: 80
```



Taints and Tolerations

<https://www.youtube.com/watch?v=mo2UrkJA7FE>

Taints - On Nodes to prevent pods to be placed on a node

Tolerations - On Pods to tolerate the taint

```
kubectl taint nodes node1 key1=value1:NoSchedule
```

<https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/>



NodeSelectors

Select the node on which a pod is to be placed

```
kubectl label nodes kubernetes-foo-node-1.c.a-robinson.internal  
disktype=ssd.
```

<https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/>




NodeAffinity

It allows you to constrain which nodes your pod is eligible to be scheduled on, based on labels on the node.

requiredDuringSchedulingIgnoredDuringExecution and
preferredDuringSchedulingIgnoredDuringExecution

<https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#node-affinity>

```
apiVersion: v1
kind: Pod
metadata:
  name: with-node-affinity
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: kubernetes.io/e2e-az-name
                operator: Inz
                values:
                  - e2e-az1
                  - e2e-az2
            preferredDuringSchedulingIgnoredDuringExecution:
              - weight: 1
                preference:
                  matchExpressions:
                    - key: another-node-label-key
                      operator: In
                      values:
                        - another-node-label-value
        containers:
          - name: with-node-affinity
            image: k8s.gcr.io/pause:2.0
```



Node Affinity - References

<https://medium.com/kokster/scheduling-in-kubernetes-part-1-node-affinity-b77c97556424>

<https://docs.openshift.com/container-platform/4.2/nodes/scheduling/nodes-scheduler-node-affinity.html>



Lab Practice

- Create pod redis with label name "type" as "db", image=redis
- Choose one of your nodes, and add a label "distype=ssd"
- Create a pod that is only scheduled on SSD nodes. (use two pods one for required and other for preferred)
- Create a service that uses an nodeport and points to a 3 pod cluster running nginx.
- List all objects in your namespace that has label "image=redis"



Lab Practice - Contd.

- Write a service that exposes nginx on a nodeport
 - a. Change it to use a cluster port
 - b. Scale the service
 - c. Change it to use an external IP
 - d. Change it to use a load balancer
- Get all worker nodes (use a selector to exclude results that have a label named 'node-role.kubernetes.io/master')

(Refer: <https://kubernetes.io/docs/reference/kubectl/cheatsheet>)



Agenda - Day 4

- **Multiple Schedulers**
- **Monitoring**
- **Metrics Server**
- **Logging**
- **Lab Practice**



Multiple Schedulers

<https://kubernetes.io/docs/tasks/extend-kubernetes/configure-multiple-schedulers/>



Monitoring

- . Node resource utilization
- . Pod container metrics
- . Application metrics



Monitoring

- . Metrics Server
- . Heapster/InfluxDB/Grafana
- . Prometheus/Grafana
- . Dynatrace
- . Datadog



Metrics Server

- . <https://kubernetes.io/docs/tasks/debug-application-cluster/resource-usage-monitoring/>
- . <https://github.com/kubernetes-sigs/metrics-server>

kubectl apply -f

<https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml>



Commands

Kubectl top node

Kubetcl top pod

<https://kubernetes.io/docs/reference/kubectl/cheatsheet/>



Logging

<https://kubernetes.io/docs/concepts/cluster-administration/logging/>

```
kubectl logs -f <container_name>
```




Lab Practice

- 1) Create a custom-scheduler called <yourname>-scheduler
- 2) Create a pod nginx which is to be scheduled using the custom scheduler created. Verify if the pod is scheduled using the custom-scheduler
- 3) Deploy metrics server by referring
<https://github.com/kubernetes-sigs/metrics-server>
- 4) Identify the node with high CPU utilization and high memory utilization
- 5) Show metrics for a given pod and sort it by 'cpu' or 'memory'



Agenda - Day 5

- **Update Deployments**
- **Commands and Arguments**
- **Environment Variables**
- **Lab Practice**



Update Deployments

Deployment Strategies:

<https://www.cncf.io/wp-content/uploads/2020/08/CNCF-Presentation-Template-K8s-Deployment.pdf>

<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>



kubectl create

kubectl get deployments

kubectl set image --record

kubectl rollout status

kubectl rollout history

kubectl rollout undo



Commands and Arguments

<https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/>

<https://phoenixnap.com/kb/docker-cmd-vs-entrypoint>



Environment Variables

<https://kubernetes.io/docs/tasks/inject-data-application/distribute-credentials-secure/#define-container-environment-variables-using-secret-data>

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap/#define-container-environment-variables-using-configmap-data>



Lab Practice

- 1) Create a busybox deployment with 2 replicas and which has "sleep 30" as arguments.. Identify the deployment strategy which is used by default
- 2) Scale the busybox deployment to 4 replicas and observe the rollout status.
- 3) Update the busybox deployment to use busybox:1.33 image and record the change so that its visible in rollout history
- 4) Undo the busybox rollout to version 2
- 5) Print the env for the busybox deployment